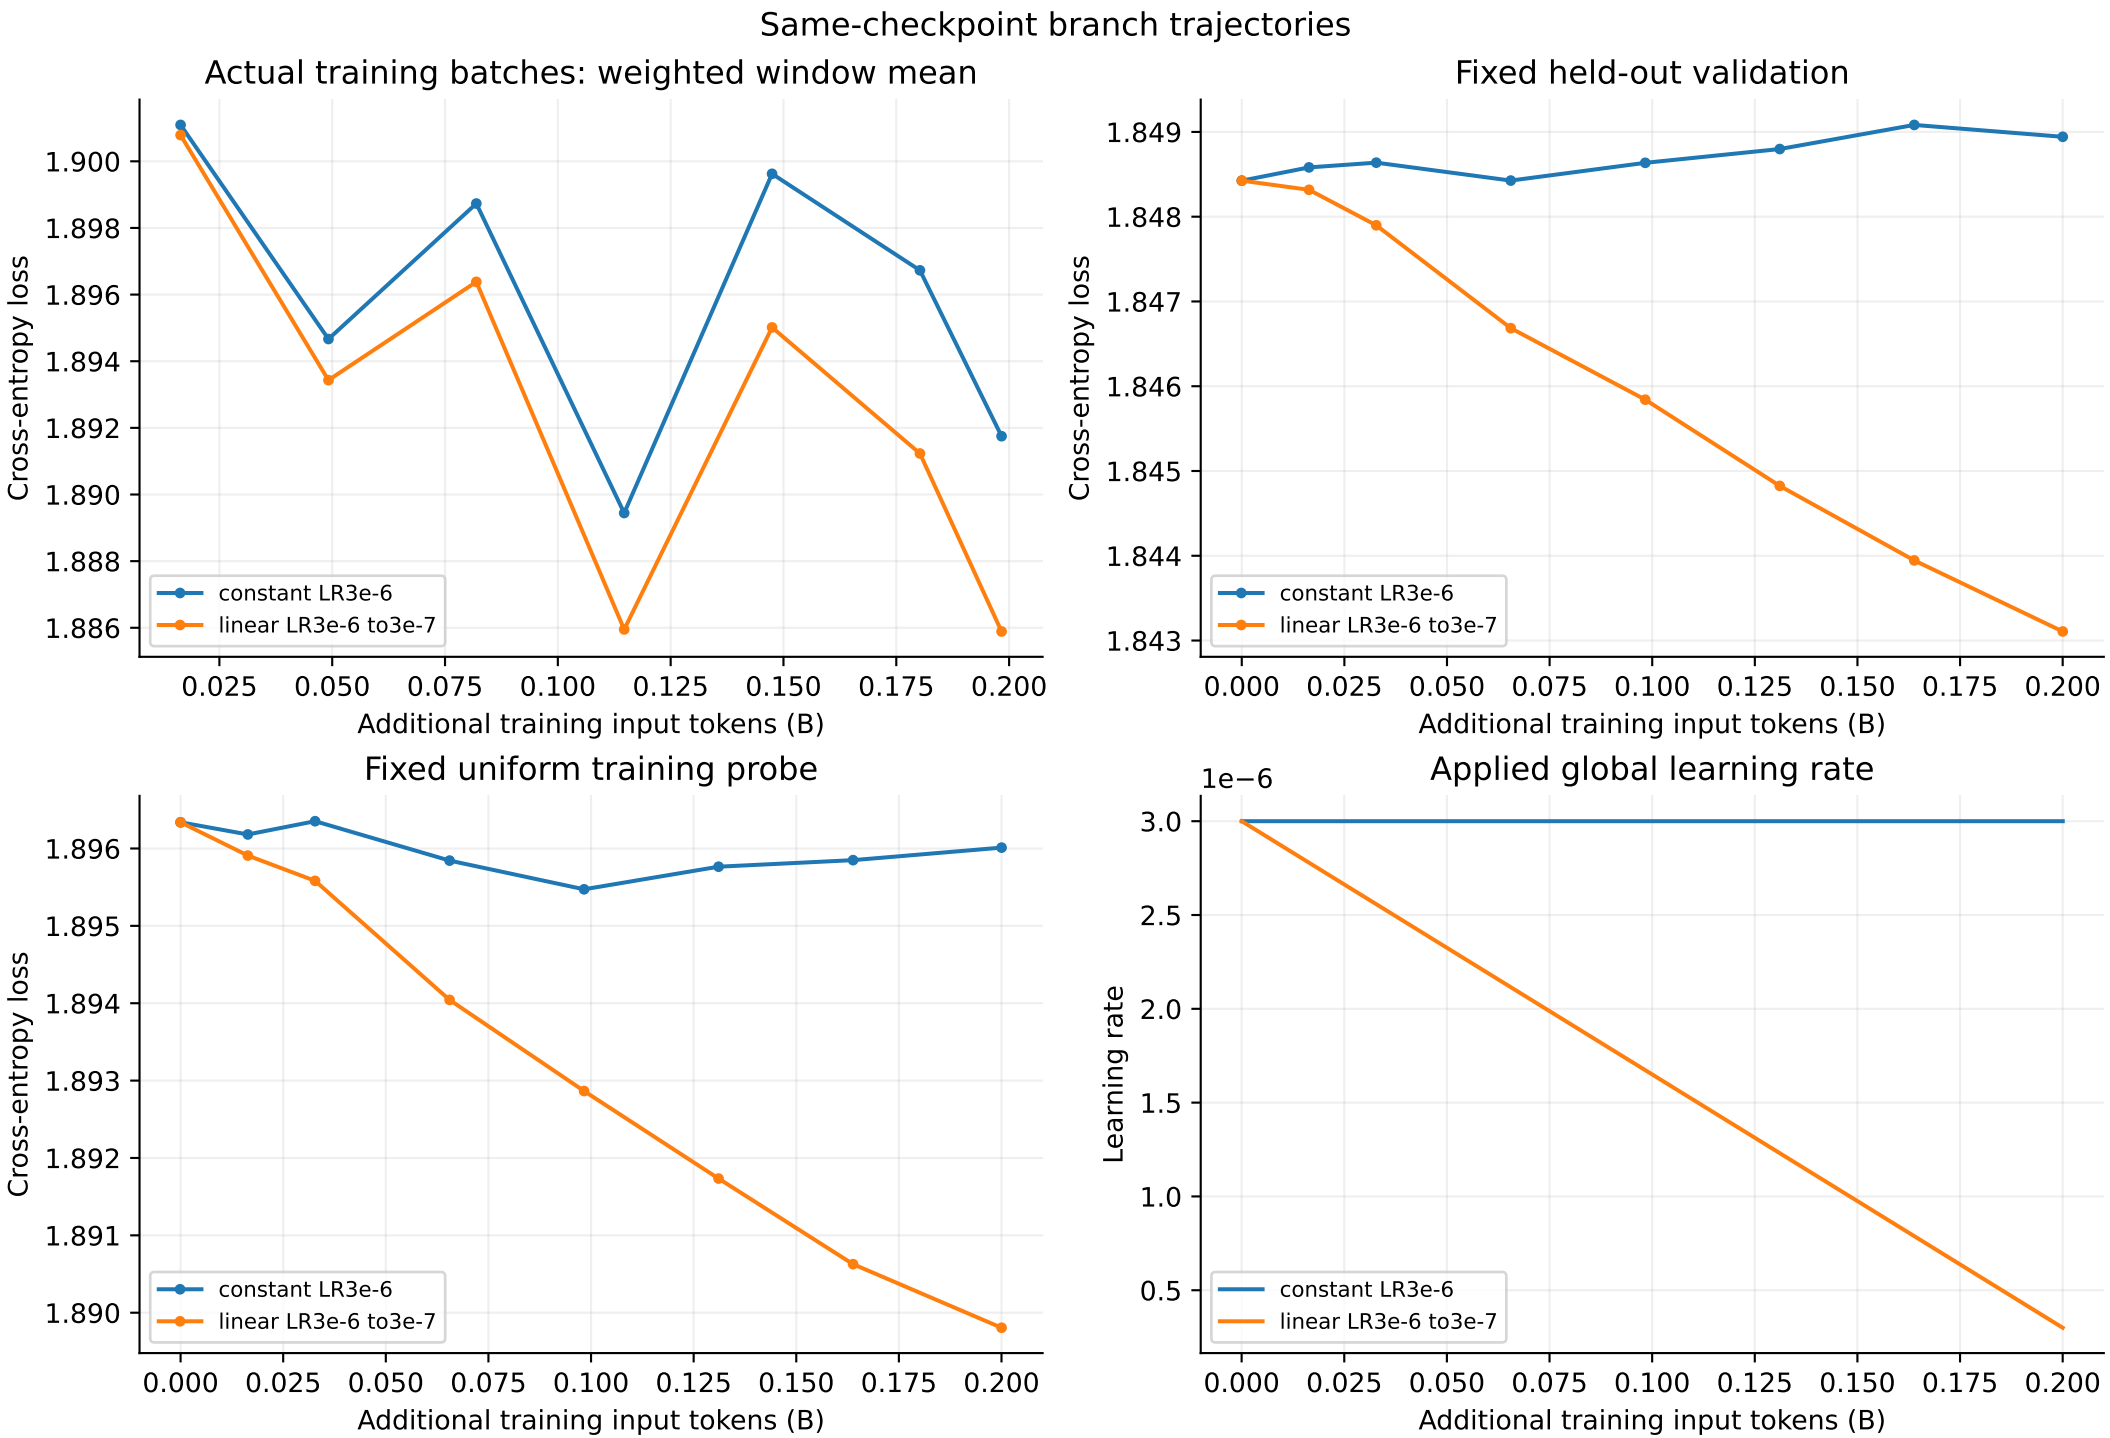
Same-checkpoint branch trajectories
Actual training batches: weighted window mean
Fixed held-out validation
Fixed uniform training probe
Applied global learning rate
Cross-entropy loss
Learning rate
1e−6
Additional training input tokens (B)
constant LR3e-6
linear LR3e-6 to3e-7